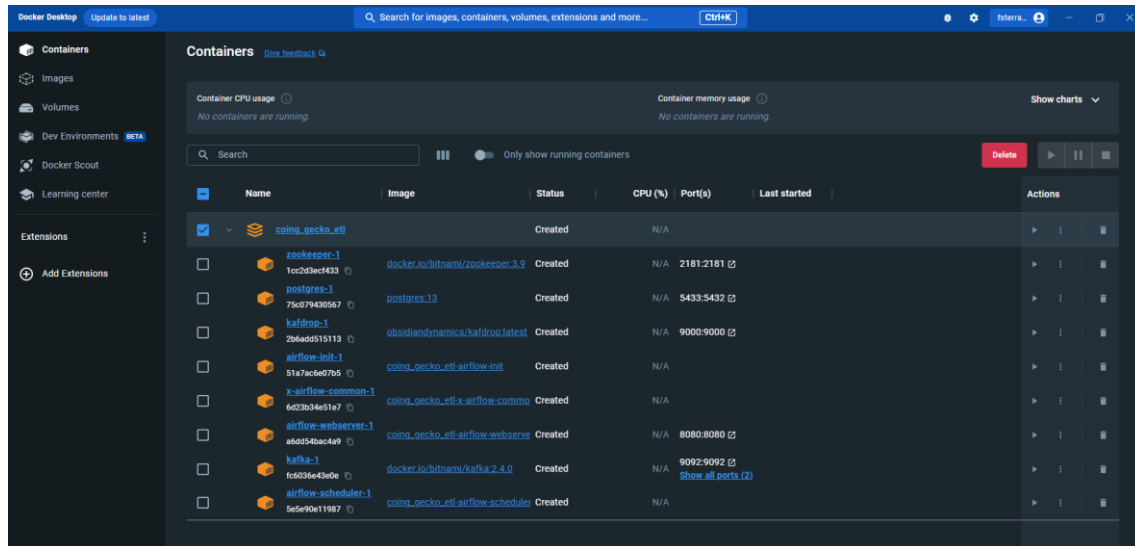


## Docker start:

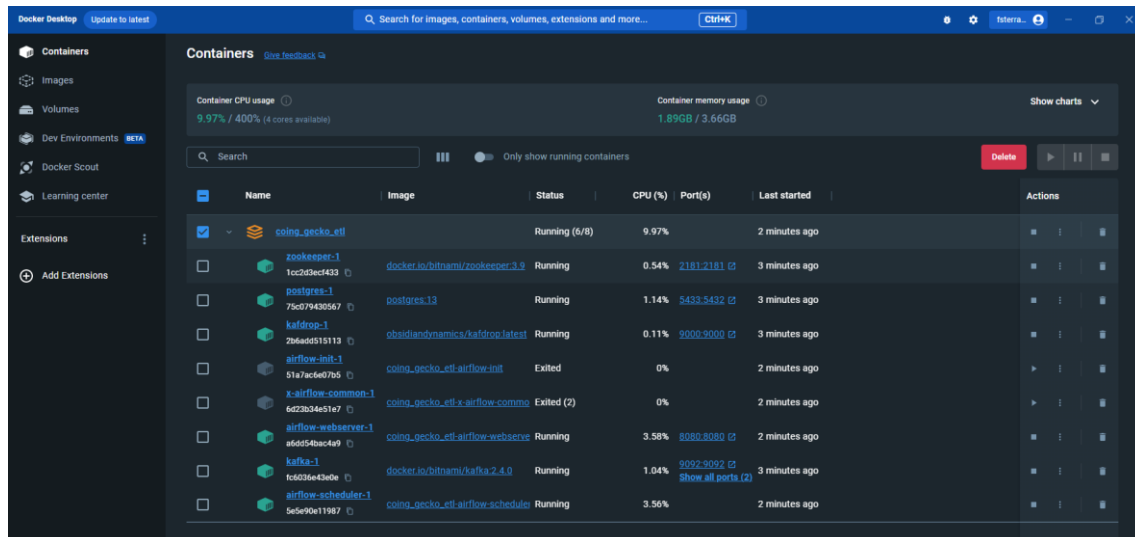
```
C:\WINDOWS\system32\cmd.exe - docker-compose up -d
Microsoft Windows [Versión 10.0.19045.4046]
(c) Microsoft Corporation. Todos los derechos reservados.

C:\Users\Franco\Desktop\Franco\Data Engineering\Kafka\Coin_gecko\coing_gecko_ETL>docker-compose up -d
```

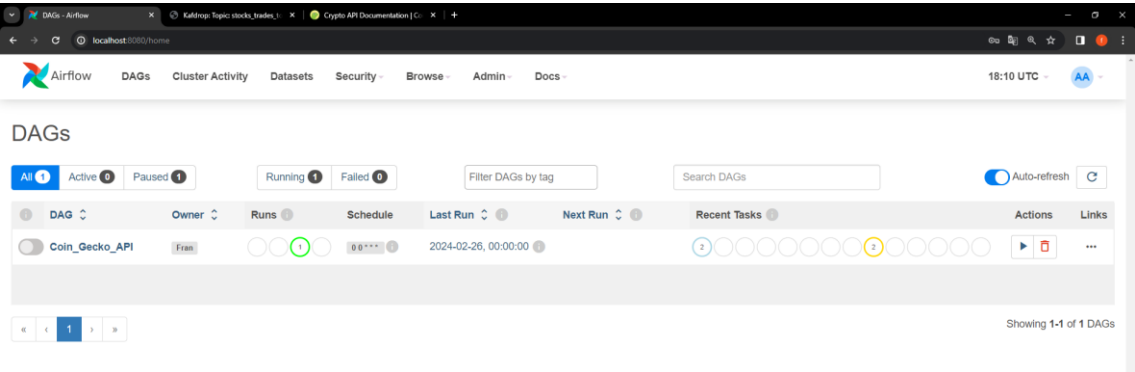
## Docker containers creation:



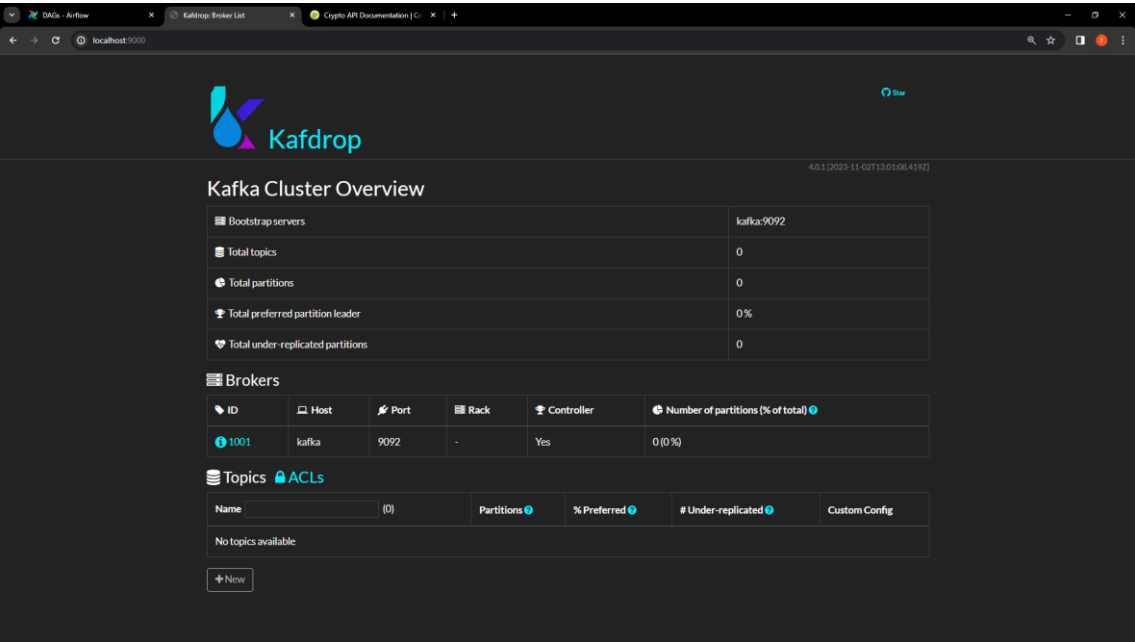
## Docker containers running:



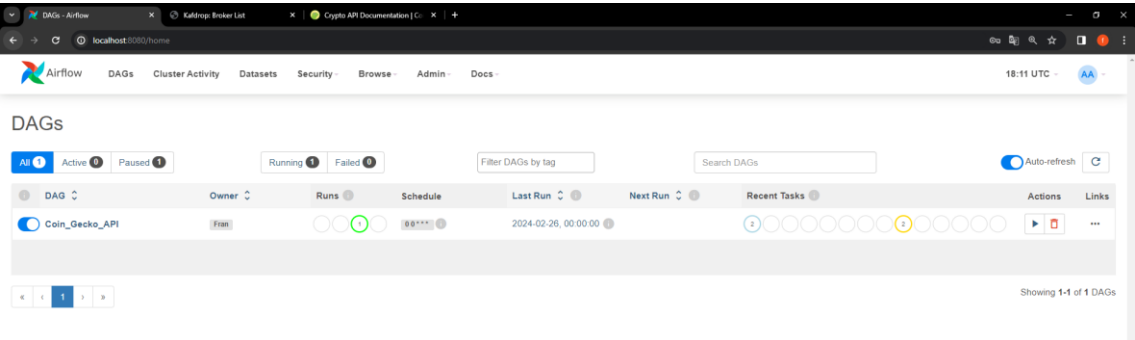
Airflow UI:



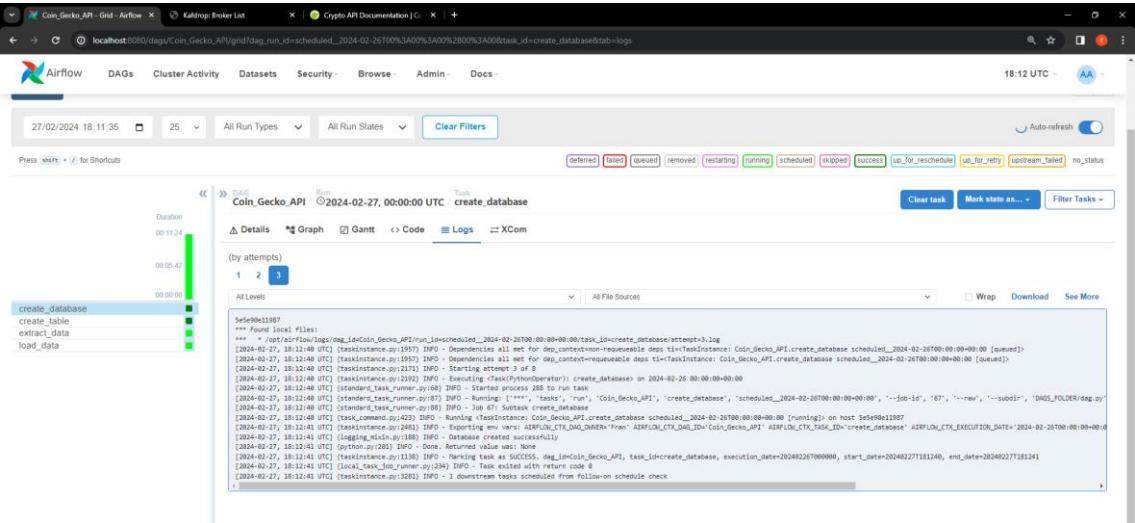
Kafdrop UI (no topics created):



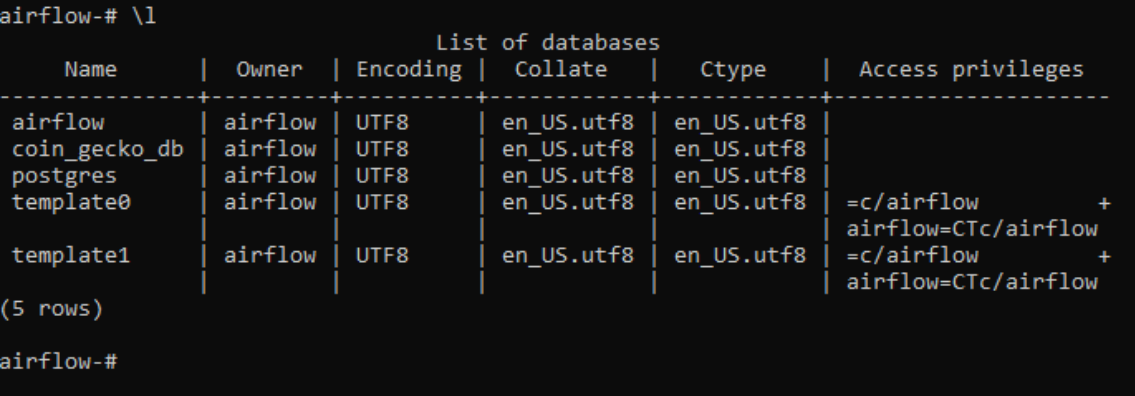
Airflow UI, DAG execution:



Database and table created in Postgres container:



Postgres Containers, PSQL CLI, database check:

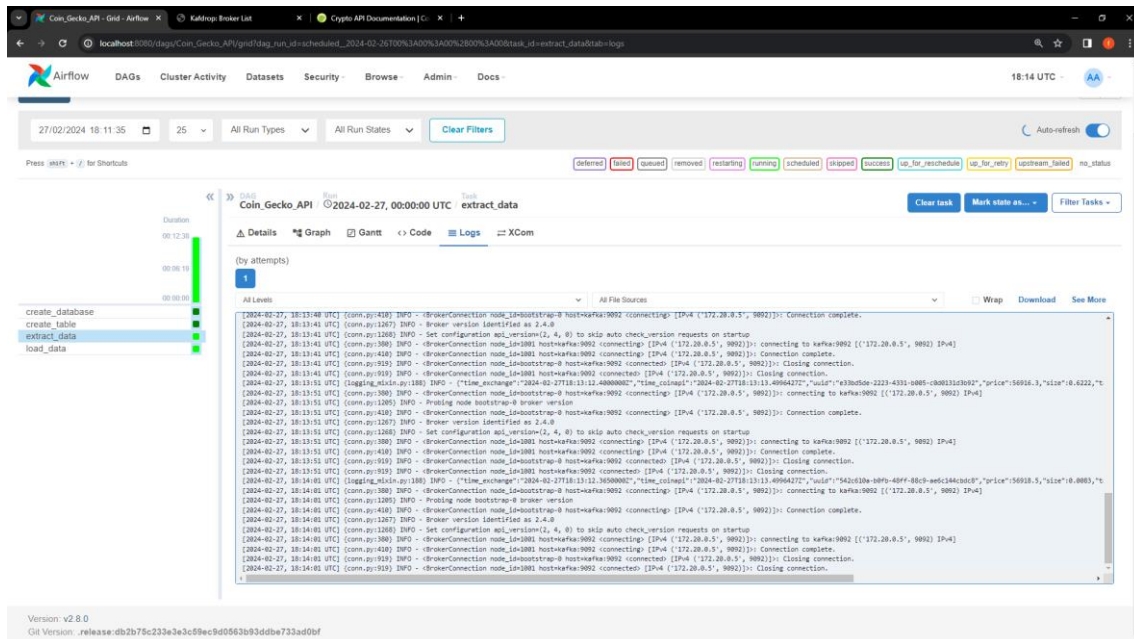


Postgres Containers, PSQL CLI, table check:

```
coin_gecko_db-# \d
List of relations
Schema |      Name      | Type   | Owner
-----+-----+-----+-----
public | trades         | table  | airflow
public | trades_id_seq   | sequence | airflow
(2 rows)

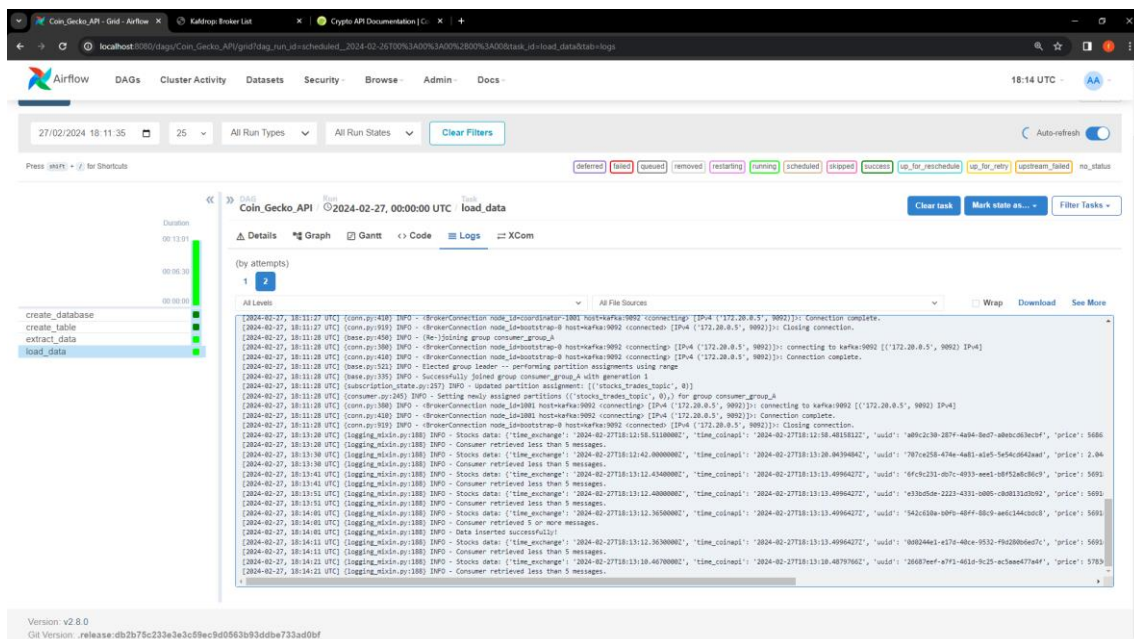
coin_gecko_db-#
```

Airflow, DB and table creation tasks successfull, Kafka producer and consumer running.



The screenshot shows the Airflow web interface for the 'Coin\_Gecko\_API' DAG. The 'extract\_data' task is highlighted, showing its progress and logs. The logs indicate that the task is running successfully, with various steps like 'create\_database', 'create\_table', 'extract\_data', and 'load\_data' being executed. The logs also show the connection to Kafka and the execution of the 'extract\_data' task.

Running consumer and loader:



The screenshot shows the Airflow web interface for the 'Coin\_Gecko\_API' DAG. The 'load\_data' task is highlighted, showing its progress and logs. The logs indicate that the task is running successfully, with various steps like 'create\_database', 'create\_table', 'extract\_data', and 'load\_data' being executed. The logs also show the connection to Kafka and the execution of the 'load\_data' task.

Kafka topic created:

Kafdrop

4.0.1 [2023-11-02T13:01:06.419Z]

### Kafka Cluster Overview

Bootstrap servers	kafka:9092
Total topics	2
Total partitions	51
Total preferred partition leader	100 %
Total under-replicated partitions	0

### Brokers

ID	Host	Port	Rack	Controller	Number of partitions (% of total)
1001	kafka	9092	-	Yes	51 (100 %)

### Topics [ACLs](#)

Name	Partitions	% Preferred	# Under-replicated	Custom Config
__consumer_offsets	50	100 %	0	Yes
stocks_trades_topic	1	100 %	0	No

+ New

Kafka topic offset increasing:

Kafdrop

### Topic: stocks\_trades\_topic

View Messages

Search Messages

Delete topic

#### Overview

# of partitions	1
Preferred replicas	100 %
Under-replicated partitions	0
Total size	9
Total available messages	9

#### Partition Detail

Partition	First Offset	Last Offset	Size	Leader Node	Replica Nodes	In-sync Replica Nodes	Offline Replica Nodes	Preferred Leader	Under-replicated
0	0	9	9	1001	1001	1001		Yes	No

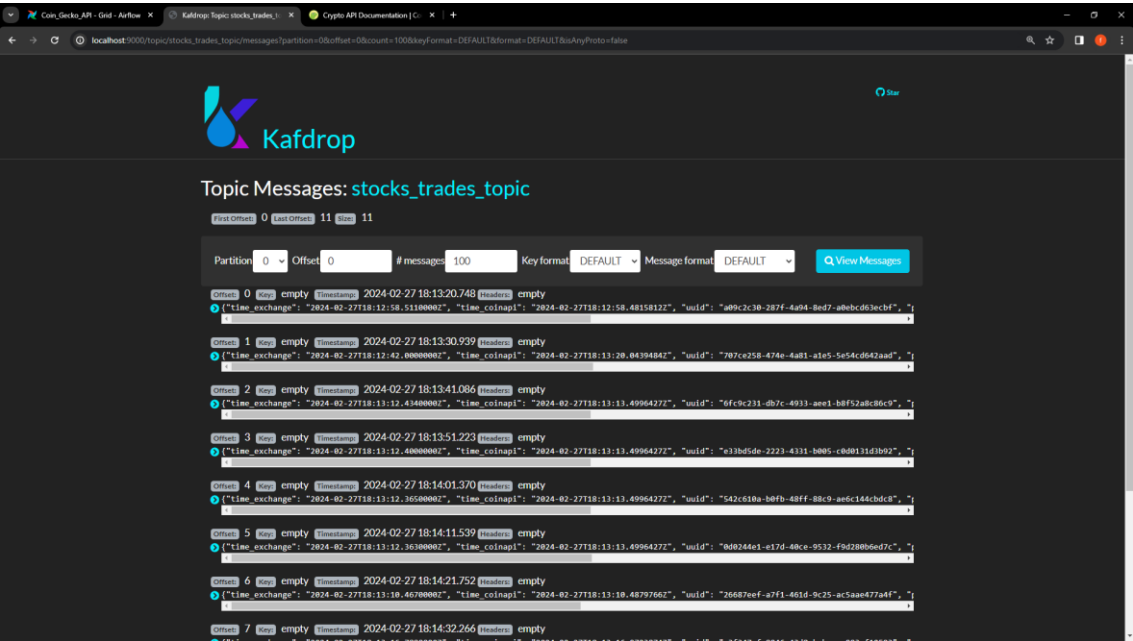
#### Configuration

No topic-specific configuration

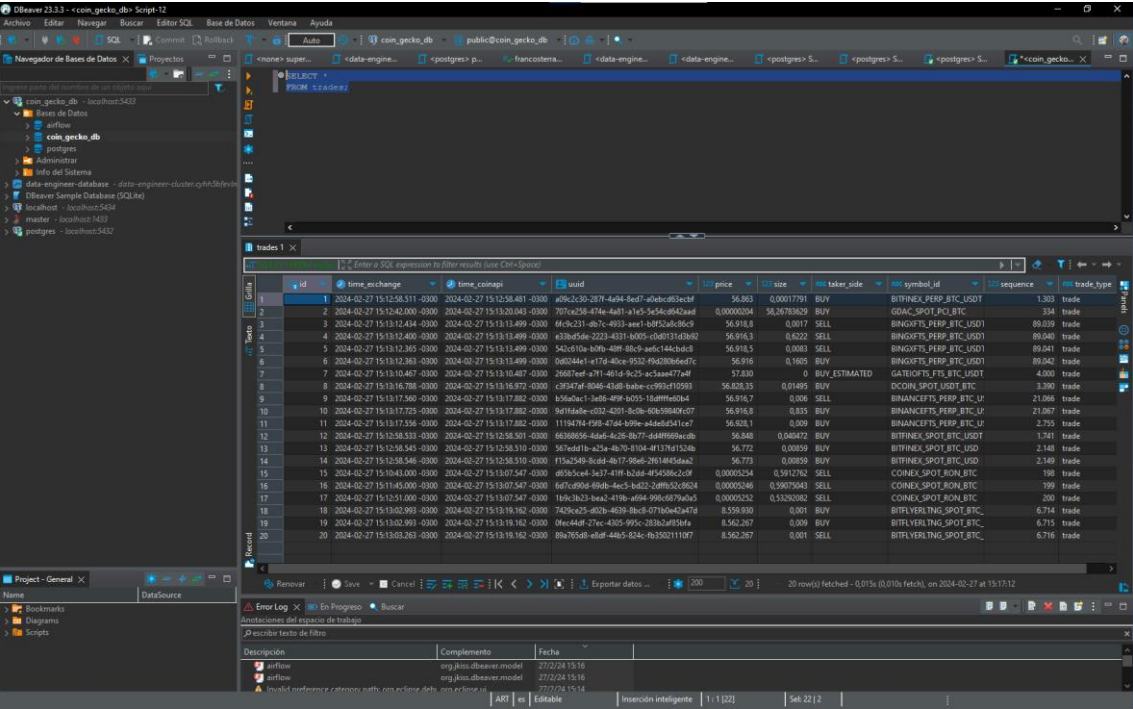
#### Consumers

Group ID	Combined Lag
consumer_group_A	0

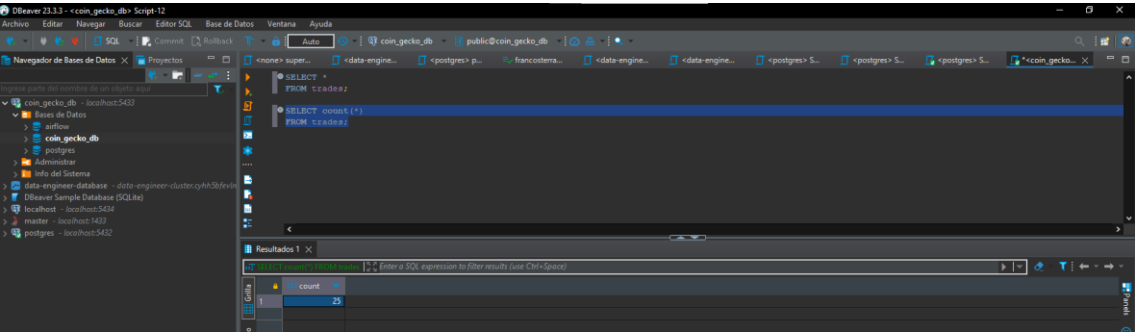
Kafka topic messages received:



Postgres DB data loaded:



Records loaded:



Records in Kafka (they are loaded every 5 records are accumulated).

